

Presenter notes – slide by slide

Rule: **the slide says it, so you don't.** These notes carry only what is NOT on the slide – the story, the reason behind each line, the window choreography, and the transition into the next slide. Point at the slide; talk about what's behind it. `build.py` bakes these into `slides/deck-presenter.pptx` as speaker notes; the organiser's `deck.pptx` stays clean.

The room (pre-poll): early responses ran ~50% “AI is part of how I run my venture” – Builder-heavy, though early responders self-select keen, so expect more mix on the day. Stance: **peer, not teacher** – “you know the tool; today is judgement and discipline.” The basics stay (the parallel watch-or-do design protects the novices); the *framing* shifts, and the notes below carry it.

Timing, honestly. The scripted material alone is ~35 minutes. The other ~25 live in the accordion: model generation time, the room running prompts in parallel (always slower than you expect), walking the floor between passes, and questions – an operator-heavy room asks more of them. If the hour runs short, the levers, in priority order for this crowd: **stretch** = expand the one-line nods into real segments (style brief → custom instructions; board of directors run live on a room idea; VET worked on a real claim; Msty demo), then the NotebookLM horizon closer (3 min, see closer-options.md). **Cut** = move 3's second pass, then the trio to one line. Never cut the two-pass reveal.

Windows: the deck, the companion site (bookmarked, `#session`), a fresh Gemini chat. Slides 6–9 are backdrops – you live in the browser there.

Tech setup (do the night before, not in the room): dedicated Chrome profile (“Demo”) – not incognito (session dies if the window closes) – logged into a burner Google account. Smoke-test it in advance: fresh accounts can hit phone-verification challenges, and first Gemini use shows consent dialogs – click through everything, run one throwaway prompt, delete that chat so the sidebar is clean. Demo in the clean account (your real account's history could steer answers via personalisation, and its sidebar stays off the projector). Fallback: the dry-run thread exported as screenshots/PDF on the desktop – don't account-switch live. Expect live outputs to differ from the dry run; the beats survive any reasonable answer. Browser zoom ~125–150% for the back row.

Slide 1 – Hook (0:00–2:30)

Slide is up as people arrive, poll QR live. Point once: *“While you arrive – one question.”*

Cold open – tell Dave's story, don't welcome anyone yet. *“Picture a sparkie – call him Dave. Sole trader out of Midland, turns over about \$300k. Somewhere in his ute is a glovebox stuffed with receipts, and he's a full BAS quarter behind. Here's the thing – Dave doesn't hate tax. He hates the chair: the sitting-down-at-9pm-with-a-laptop part. He'll do a 12-hour day in 40-degree heat, no worries; ask him to file a receipt and he'll find literally anything else to do.”*

Then the turn: *“In about ten minutes, we're going to pitch an app that rescues Dave to an AI – live – and you're going to watch that AI lie to us. Then we're going to fix it.”*

Then use the poll as the thesis, demonstrated. Glance at the histogram and say what it shows: *“Half of you told us AI is already part of how you run your venture. Good – because it's also part of how your competitor runs theirs. When everyone has the brilliant assistant, nobody has an advantage. Today is*

about where the advantage moved." This one beat converts the advanced half from "I know this already" to "this is about me."

Now the provocation – deliver it as your own claim, don't read the quote block. The question at the end ("where does your edge actually live?") is the whole hour; let it hang for two seconds.

Show of hands (who's used AI this week) – count it, don't discuss it. Housekeeping: slides are signposts, the site is the handout, scan now or later.

Do: if the live histogram contradicts the pre-poll (more A/B than expected), drop the "half of you" line to "*many of you*" and slow the follow-along instruction later – the beat still works.

Advance: off the hanging question.

Slide 2 – The thesis (2:30-4:00)

Slide 1 was the competition angle (everyone has it). This slide is the capability angle – don't restate, change the question. Ask the room: "*Why can't I just ask AI to write me the next viral app, or the next bestseller? It has read every bestseller ever written. It knows every app that ever went viral. So – why can't it?*" Let it hang. If someone offers an answer, take it and build on it, don't correct it.

The answer plants the hour's key word: it gives you the **average**. It's a machine for the centre of the distribution – the most plausible next thing, synthesised from everything everyone has already done. Fluent, usually right, and identical for whoever asks. But a bestseller lives in the *tail*: a taste call, a timing bet, a contrarian read the average disagrees with. The tail isn't in the model. It's in you.

Now point at the slide's last row and give the tail its name: **judgement** – taste, context, the variation no identical model can reproduce. "The average" comes back twice, on purpose: the generic pitch in move 3 is the average, live – and the trust tool's first axis is literally named for it.

Advance: the moment "the tail is in you" lands. Don't linger – this slide is a promise, not a proof.

Slide 3 – Walk the talk (4:00-4:45)

The slide shows the résumé; you supply the tone – dry, fast, almost throwaway. The unspoken question you're answering: "has this person actually done it, or do they just lecture about it?"

The one line to deliver with a beat of silence after: "*the deck you're watching, the site you scanned, and the poll you just answered were all built in conversation with AI.*" Don't elaborate. Understatement is the credibility.

Then set expectations – the credibility earns this refusal: "*So yes – you can build real software in conversation with AI; that's where those thirty tools came from. But I'm not here to show you how clever I am with AI – and in one hour I couldn't make you clever with it either. The tool is the easy part, and everyone in this room has the same one. I'm here for the other part: the judgement you bring to the conversation. That's the only part that's yours.*" Keep it this short – longer reads as apology. (This also pre-plants the literacy line for the close.)

Transition: "*So – Dave's app. Let's find out what it's up against first.*"

Advance: on that line.

Slide 4 – Trap 1: Sycophancy (4:45-8:00)

The slide names the trap; you supply the *why* and the cost. Why it happens: the model was trained on human approval – agreeable answers score better. It's not lying, it's people-pleasing at scale.

Frame both traps for heavy users, not beginners: *"These two get worse the more you use it – familiarity breeds trust. If AI already runs half your workflow, you're consuming more of its output than anyone, and this tax compounds."* The advanced half thinks they're past this; that belief is the trap.

The cost, for founders specifically: praise stops iteration. *"The most dangerous thing a founder can hear in month one is 'great idea' – because you stop stress-testing exactly when stress-testing is cheapest."*

Tease, don't teach, the defence: *"You can prompt straight past this – you'll watch it happen in about twenty minutes."*

Advance: on the tease.

Slide 5 – Trap 2: Gell-Mann amnesia (8:00-11:30)

Tell the origin story – it's not on the slide and it's why the name sticks. *"Michael Crichton named this after his friend, the physicist Murray Gell-Mann. You open the newspaper, read an article on something you actually know, and it's riddled with errors. Then you turn the page to the finance section – and read it as if it's gospel. Same paper. Same journalists. You knew they got your field wrong, and you forgot it one page later."*

Then map it: your domain = the trade, the product. The page you turn to = the market size, the regs, the numbers. AI writes every page with the same confident fluency.

For the daily users specifically: *"This one bites hardest the more you delegate – because the more of your venture AI runs, the more pages you're reading that you can't check."* The heavier the usage, the more relevant, not less.

Defence in one line: be *most* sceptical where you know *least* – inverted from what feels natural.

(Dismissal Fallacy is deliberately cut – don't mention it.)

Transition: *"Right. Everyone open a chat window – we're going to Dave's rescue."* This is the minute-12 mark.

Advance: as laptops open.

Slide 6 – LIVE: one idea, three moves (12:00)

Choreography slide – almost everything here is Do, not Say.

Say the IP beat first, before anyone pastes (verbatim, it's rehearsed): free chat is someone else's cloud; sensitive idea → work the tradie example now, run yours tonight in Msty; frontier power vs privacy is a founder decision that never goes away. This gives permission to watch instead of paste – some of the room needs that permission.

Do, in order: (1) companion site on the projector, scroll to the session prompts – say *“everything I paste today comes from your handout, in order – nothing up my sleeve.”* (2) Fresh chat tab beside it. (3) Confirm the room is in: *“One of the three free chats on screen – you run your idea in parallel, or just watch. Both work.”*

Mixed models on the floor – name it as lesson one, then set one rule. (Dry runs: Gemini writes a market report, Claude writes one and asks for your wedge, Fable offers a 10-minute Research mode.) Say: *“You’re on different models, and you’ll get different answers – different formats, different voices; one of them may offer to go away and research for ten minutes. That’s not a bug – it’s the first finding: none of them will be specific to you until you make them.”* The rule: decline research/agent modes today – *“we’re conversing, not commissioning. That button is ten minutes of very thorough research into the market everyone can see – press it in move 2, when you’ve got an assumption worth verifying.”* (Research mode = the verify column of the trust tool – the callback lands at slide 11.)

Harvest the differences, don’t fight them: after each pass ask *“who got something different?”* – take one or two, max. A clarifying question, a research offer, a different format: free teaching material. More than two and you lose the clock.

Fast-finisher escalation (operator-heavy room will need it): *“Done before me? Go to the starter prompts on the site and run the board of directors on your idea.”* Say it once here; repeat as needed during the moves.

Advance: to slide 7 as you copy prompt 1. From here to slide 10 the deck is a backdrop.

Slide 7 – Move 1: Validate (12:00-24:00)

Pass 1 – frame it as the control, not a lesson. An operator will bristle at *“here’s how you might prompt”* (*I’d never prompt like that*). So say it straight: *“None of you would paste a prompt this lazy. But this is the baseline – watch what the market gets.”*

Run the live demo on Gemini. Dry runs (18 Aug) on both Gemini and Claude: Gemini obeys the Format line – three assumptions, one line each – which gives the clean essay-vs-three-lines contrast on screen and leaves prompts 3-4 something to do. Claude’s answers are richer but front-run move 2 and the cheapest-experiment beat (it volunteered a \$300 shoebox test in pass 2). Claude output = Q&A/backup depth, and a bonus line: its naive answer literally asked for the wedge – *“even the model knows what’s missing from this prompt: it’s asking for my edge.”*

Don’t promise flattery – promise sameness. Both models returned no love-fest on the naive prompt but a competent, critical consultant’s report – competitors, risks, a validation plan. Sycophancy still shows on opinion-asks (*“what do you think?”*), not task-asks. The reveal works for any outcome: *“I ran this on two different companies’ models last night – different vendors, same report. Notice the one thing this answer can’t be: specific to me. If it flatters me – everyone gets flattered. If it’s this market report – everyone gets this exact report. Either way: the average, live.”* (If it DOES gush: name trap 1 from slide four, live – a gift, take it.)

If a confident competitor list appears (Claude named six, several unverifiable): that’s trap 2, live. *“Six competitors, named with confidence. Which are real? You can’t tell from up here – smoothness isn’t accuracy. That’s why VET exists.”* Assert unverifiability, never that a specific one is fake – you can’t check from the podium either; that’s the point.

Pass 2 – the point is strategy, not prompt technique. Half this room knows RTCF mechanically; almost none of them treat their context as the moat. Before sending, point at the slide's Role · Task · **Context** · Format and land it at that altitude: *"Context is the only line the model can't invent – and the only line your competitor can't copy. It's your defensible asset; the discipline is injecting it every single time."* Dave's context – the \$150-400k sole trader, the bookkeeper-shaped competitors, the penalties-not-deductions bet – took a founder to know.

When it returns, show the aiming – the three assumptions are the three context lines handed back as testable risks, one for one (dry run confirms Gemini does this): servo photo → will they open an app at the pump; bookkeeper bet → standalone value vs churn back to full accounting; penalties bet → will saved fines be attributed and paid for. Then the reveal, slowly: *"Same model. Same idea. Sixty seconds apart. The AI did not get smarter – I aimed it."* Thesis, second landing.

Hold in your head: this same model calls the idea a PIVOT with a \$50 experiment in move 2. Don't foreshadow it to the room – the surprise is the payoff.

The suggestion chips (Gemini offers "where should we focus next?" after pass 2) – half a beat, not an apology: *"See these? The model's offering to drive. They're reasonable – and they're follow-ups to its answer, not my priorities. At home, click them when they match where you were already going. Just notice who decided. Today, we drive: why does this fail?"* (If its first chip is close to the pressure-test anyway, add: *"its suggestion is near where I'm headed – but I'll frame it my way, as an attack."*) Same discipline as declining the Research button – direct it, don't delegate applies to navigation too.

Advance: to slide 8 as you copy prompt 3.

Slide 8 – Move 2: Pressure-test (24:00-36:00)

Paste prompt 3. Read the three failure reasons out loud – slower than feels natural; this is the room hearing an AI be genuinely useful for the first time today.

The load-bearing assumption is the beat (dry run: Gemini names the injected contrarian bet, verbatim, as the thing most likely to be wrong). Read that final sentence aloud, slowly, then: *"In prompt 2 I told it my contrarian bet. In prompt 3 it's telling me my bet is the part most likely to be wrong. It's not attacking the idea – it's attacking my claim. It can only do that because I gave it the claim."*

Optional one-liner if it fits: context compounds – aimed once in prompt 2, still aimed two prompts later. Then prompt 4.

Scale/Pivot/Kill: gloss, don't teach (this room knows the words): *"Scale – double down. Pivot – change the bet. Kill – stop."* Five words each as you paste prompt 4; the pills on the slide do the rest.

The arc – name it the moment the verdict lands (dry run: Gemini returns a bold PIVOT, calls the app-first plan "a dead end", and prescribes a \$50, 5-day Wizard-of-Oz WhatsApp test): *"Prompt 1 gave me the market report everyone gets. Prompt 4 just told me my plan is – read it – a dead end. Generic, to aimed, to a verdict – that arc is the entire session in one thread."* Two more points at the screen: the experiment carries a built-in kill criterion (no photos to WhatsApp in 14 days → KILL), and the test itself is aimed – it attacks the habit gap using the thread's own context (zero friction, no new app, where their thumbs already are).

Normalise the kill – the room needs to hear a lecturer say it, and the kill criterion is the moment: *"It just handed me a way to find out for fifty dollars whether to stop. A fast cheap kill is the best deal in startups;*

the expensive version of this lesson costs eighteen months."

Name the tool as you leave: what we just ran is the *devil's advocate* – the first of three; the other two get one line each later.

Time check at 36: behind → cut move 3's pass 2, never the reveal.

Advance: to slide 9 as you copy prompt 5.

Slide 9 – Move 3: Sharpen the pitch (36:00–45:00)

First: accept the pivot, out loud – the human decides. *"The verdict said pivot. I accept it – my call, not its."* The site carries BOTH versions of prompt 5 (pivoted = the paste-as-is; original = the labelled alternative): *"Your handout has two versions of this one – the original idea, and the idea as it now stands after the pivot. I'm taking the pivot version, because I accepted the verdict. Take whichever matches your call. The prompts are the method; the idea is always yours."* Room instruction: *"pitch YOUR idea as it now stands – Scale, pitch it proudly; Pivot, pitch the new shape; Kill – pitch your next idea."*

Pass 1 (prompt 5): when the generic pitch appears, ask the room, don't tell: *"Would you remember this tomorrow? Neither would an investor."* Then cash slide 2's promise: *"This is the average, live – the centre of the distribution. Every founder who asks gets this pitch."*

Pass 2 (prompt 6 – take the pivot version, matching your prompt 5 choice; the original stays as the site's alternative): before sending, read Dave's voice sample aloud with the accent it deserves – "they don't hate tax, they hate the chair" gets a laugh and *earns* the contrast. The room pastes their own writing via the template. Bonus if it fits: Dave's sample was already pitching the pivot – *"one photo at the servo, snap it, chuck it, drive off"* IS the WhatsApp pipeline. *"His own words found it before the model did."*

Put both pitches side by side – they're two consecutive replies in the same thread; copy the two 60-words into a blank doc (or scroll between them) so the room sees both at once – and shut up for five seconds.

The comparison beat (dry-run confirmed: pass 1 comes out edge-aware but in agency voice – use the "could be anyone's" caption): point at the two OPENING lines. Pass 1 opens with an instruction ("Stop stuffing paper receipts into your glovebox") – an ad; any copywriter writes it. Pass 2 opens with a belief ("Tradies don't hate tax – they hate sitting at a laptop at 9pm") – then the customer, then the bet. *"The first pitch sells a product. The second sells a founder's view of the world – and investors fund the second, because the product will change and the view is what steers it. Nothing changed except whose words went in."*

Then the meta-beat – once, here, nowhere else: *"And notice what just happened in the silence. Nobody told you the second one was better. No model ranked them. You judged it – instantly. That comparison you just made in your head is the one thing that never went to the cloud. That's the human in the loop – and it's the part you can't delegate."* The room just proved the thesis on themselves; caption it and move on.

One-line nod: the style-brief trick makes this permanent – it's linked on the site. **(Stretch lever for an advanced room: make this a 3-minute segment – generate a style brief live from the voice sample and paste it into custom instructions. Daily users use AI constantly and almost never systematise this.)**

Advance: to slide 10, energy up – the live block is done, stick the landing.

Slide 10 – Push back: three ways (45:00, ~1 min)

The slide lists them; you only connect and locate. Devil's advocate – *"you just watched it."* Board of directors – one colour line: *"the CFO in that room will hate your pricing; better to hear it from a fake CFO for free."* VET – point forward: *"that one earns its keep in two minutes."* All three prompts on the site.

Planned stretch #1 lives here (likely needed – an operator room runs the live block fast): if you're at this slide before ~45:00, run the board of directors live in the same thread, using the adapted prompt (site version stays generic): *"Answer as a panel, each in their own voice: (1) a cautious CFO, (2) my most skeptical customer, (3) a growth-hungry marketer. Each gives an honest take on this WhatsApp receipt photo idea, then stops. I will decide."* Point at your own last line before sending: *"the whole hour in five words."* 3–4 minutes, and it's the beat the organisers loved.

Three beats in the output (dry-run confirmed): (1) they disagree – CFO "not venture-backable", customer "I'll stick to the glovebox", marketer "pure gold" – *"an answer machine gives you one answer; a board gives you a decision to make."* (2) The cheerleader took a seat: the marketer IS trap 1, relocated – you weigh that voice, not obey it. (3) The customer surfaces failure modes nothing else found (grease, dropout at the trade desk, faded thermal paper).

VET: discuss by pointing at this very output – no extra prompt. The board just made three smooth, specific, unverified claims (per-message API costs, thermal fade, "\$100k ARR on referrals"). *"Which do you VET first? The one your next decision leans on – the API cost, inside the CFO's kill argument. Verify the source, Explain it back, Test the edges. Tonight's homework; prompts on the site."*

Advance: briskly. This slide is a landing, not a lecture.

Slide 11 – The trust tool (45:00–53:00)

Open with the callback – it makes the grid click instantly: *"Remember the question from the start – why can't AI write the next bestseller? Because it gives you the average. That's this axis. Where the average is good enough – drafts, scans, brainstorm – lean in. Where you need the tail – the call, the number, the bet – that's yours."* The grid stops being a framework and becomes the answer to the opening question.

The grid is on the slide; you supply the sorting exercise. Give the room two concrete calls to make: *"Market-sizing first pass – which box?"* (lean in, verify lightly). *"The revenue number in your pitch deck?"* (human owns it, full stop). Let them answer out loud; correcting a wrong answer teaches the grid better than presenting it.

The keeper line, verbatim, once: *"Where it hurts to be wrong, you stay in charge."*

Close the loop from slide 10: VET is *how* you do the verify column – it's in the starter prompts as "VET the claim."

Advance: at ~53:00.

Slide 12 – The method, in three sentences (53:00-58:00)

Everything is on the slide. Your only job: connect each line to the moment they watched it happen – context → pass 2 of move 1; past the flattery → move 2's verdict; sceptical where you know least → the Gell-Mann page-turn. *"You didn't learn a theory today. You watched all three."*

Thesis, third landing: thinking partner, not stand-in.

Advance: into the close.

Slide 13 – Takeaways (58:00-60:00)

Make it one action, not four bullets: *"Tonight, before you lose the nerve: run move 1 on your real idea. Naive first, then with your edge. When you feel the difference – that feeling is the method."* Everything else – the builder, the library, Notebook, Msty – is on the site, and the site is bookmarked.

The promise (retention hook – share the LIVE thread, not the dry run): *"Everything I typed today – the full conversation, and the notebook built from it – will be on the companion site this afternoon."* Share from the burner account (test tonight that both share links open for non-owners; notebook shares may need sign-in). The dry-run versions are the fallback only if the live beat got cut or broke. Add the links to the site after the session.

Time-permitting horizon closer (closer-options Option B, upgraded – no slide needed, slide 13 stays up; ALWAYS before the triad, never after): save the live conversation to a Gemini Notebook and generate the ONE artifact you pre-tested tonight for speed (text artifacts fast; audio overview takes minutes – if you want it, kick it off and keep talking). Framing, so it's thesis not showcase: *"Notice what this is not – it's not the average. It's grounded in OUR conversation; it cites what we said. Same model family, different placement on the grid."* Chain the research use off the morning: *"Remember the Research button we didn't press? Tonight, that plus a Notebook is how you VET a competitor list – everything public, grounded, cited."* One sentence on the ecosystem, no demo: *"Gemini's also growing agent features and plugs into Docs and Sheets – the honest trade: not the strongest model, but when your method is conversation, the model race matters less than workflow fit and the price of free."* Hard cap 3 minutes.

Q&A deflection, whenever a tool question comes ("which one should I use?", "what about Grok Bot?"): *"The names change monthly. AI literacy isn't about the tool – it's about how you converse with it. Everything you watched today works in any of them."* One sentence, back to the method.

Close on the line, then stop talking: *"Direct it. Don't delegate. Never surrender."*

Do: leave the slide up as they pack – the QR keeps working.